

# Arjun Gullbadhar

AI Engineer — production LLM systems & evaluation

Delhi NCR, India (open to relocation anywhere) | [arjunggullbadhar01@gmail.com](mailto:arjunggullbadhar01@gmail.com) | +91 88476 27161 · +91 97797 55748  
[arjunggullbadhar.com](http://arjunggullbadhar.com) | [linkedin.com/in/arjun-gullbadhar-11600b167](https://linkedin.com/in/arjun-gullbadhar-11600b167)

---

## PROFESSIONAL SUMMARY

---

AI engineer with five years across data, backend, and applied-LLM systems, now specializing in **production LLM systems and their evaluation**. At PharmaEdge, builds the evaluation, retrieval, vision, and research pipelines behind a pharma competitive-intelligence product. Previously built a 41-tool production MCP server for an SEO platform and shipped 10 consumer iOS apps solo. Conference speaker on LLM testing in production.

## EXPERIENCE

---

**AI Engineer, PharmaEdge.ai** (pharmaceutical competitive intelligence) Nov 2025 – Present

- **Built the flagship news-intelligence system end to end**, from ingestion to client-ranked delivery: pharma news lands every 30 minutes, is clustered into stories, and is ranked per client against tracked drugs, competitors, and therapeutic areas; raised story-clustering F1 from 0.62 to 0.855 on a hand-labeled golden set by combining embedding similarity, LLM-extracted entities, and entity-overlap scoring with a news-type taxonomy.
- **Built the conference-intelligence analyzer end to end** — a multimodal pipeline that turns conference posters, booth photos, and slide decks into structured competitive intelligence, routing each asset through adaptive OCR and vision paths (GPT-4.1/4o) and digitizing Kaplan–Meier survival curves into typed data; cut per-asset processing time 12–42× across asset types.
- **Established the evaluation layer for every LLM service**: versioned golden datasets with per-merge regression and cost gates that block any prompt or model change that regresses quality or raises cost, LLM-as-judge review calibrated against human labels, and drift detection; drove recurring output-formatting failures to zero with instructor + Pydantic structured outputs.
- **Built the investor-relations call intelligence system**: chunked Whisper ASR transcribes live earnings calls, biased by a domain vocabulary mined from each company’s previous three quarters of calls; client-tuned watchlists flag relevant moments to the right client as the call happens, and every call ships with a generated prep brief of questions and updates to listen for, plus summaries and extracted pipeline milestones.
- **Built a multi-stage deep-research workflow and RAG chat** over the proprietary corpus — pgvector retrieval with reciprocal-rank fusion and LLM reranking, plus retrieval-verification and citation checks so generated claims trace back to sources; structured extraction over ClinicalTrials.gov feeds the corpus typed trial records.
- **Built the observability pipeline for the LLM services**, logging every model call with tokens, latency, cost, and status, and tracking quality and spend across services.

**AI Engineer (Contract), Welzin.ai** (SEO content platform) Aug 2025 – Oct 2025

- **Designed and built a production MCP server exposing 41 tools** for content generation, scheduling, social publishing, and Search Console insights, used live by the client’s agent workflows; event-ID deduplication ensures retries never publish duplicate social posts.
- **Delivered three client-facing workflows end to end** across NestJS and Python services: a SERP-grounded SEO wizard, a server-rendered carousel studio, and a LinkedIn-style social composer.

**Indie iOS Developer & Technical Writer***Mar 2024 – Nov 2025*

- **Shipped 10 consumer iOS apps solo in Flutter**, owning product design, LLM backends, App Store review, localization for 15+ storefront markets, and RevenueCat subscription billing; several were LLM-powered, including a calorie tracker built on chained model calls.
- **Curated a domain training dataset and fine-tuned Qwen** (Unsloth SFT) for domain-specific generation.

**Data Analyst, Zebra Technologies** (Antuit.ai)*May 2021 – Jan 2024*

- **Built pricing-analytics models for GlaxoSmithKline** covering 50+ markets; outputs fed directly into annual pricing strategy reviews with client stakeholders.
- **Developed demand-forecasting solutions** for retail brands including J.Crew and Madewell, informing inventory and promotion decisions.
- **Designed and operated PySpark ETL pipelines on Azure Databricks** processing 10M+ transactions daily.

**TALKS & WRITING**

---

- **Speaker, AAIF & CNCG Summit 2026**: “Prompt Regression Testing using Argo”
- **Speaker, OpenSearch Delhi Meetup 2026**: “Make Sense of a Pile of Stuff: Embeddings Beyond Search”
- **Medium Top Writer in Artificial Intelligence** (awarded Dec 2022) — 250K+ cumulative readers on the technical blog, [medium.com/@arjungullbadhar](https://medium.com/@arjungullbadhar)

**SKILLS**

---

**LLM & agents**: OpenAI, Anthropic (Claude), and Gemini APIs; structured outputs (instructor, Pydantic, PydanticAI); LangGraph, LangChain, LlamaIndex; MCP server design; multi-agent orchestration; LoRA/QLoRA fine-tuning

**Evaluation & observability**: golden-set regression testing, per-merge quality and cost gates, LLM-as-judge calibration, RAGAS, drift detection, LangSmith/Langfuse, full model-call logging

**Retrieval & multimodal**: PostgreSQL/pgvector, hybrid search, reciprocal-rank fusion, LLM reranking, OCR/vision extraction, Whisper ASR

**Engineering & data**: Python (async, typing), SQL, Bash; FastAPI, Docker, GitLab CI, S3, SSE streaming; pandas, scikit-learn, PySpark, Azure Databricks; Flutter (iOS, full App Store lifecycle)

**EDUCATION & AWARDS**

---

**B.E. Computer Science & Engineering** (Cloud Computing with IBM)*2016 – 2020*

Chandigarh University

Kaggle Elo Merchant top 19% global · Hackoverflow Hackathon 2nd of 108 teams · TechGig ML Code Gladiators national finalist