

Arjun Gullbadhar

AI Engineer — production LLM systems & evaluation

Delhi NCR, India (open to relocation anywhere) | arjunggullbadhar01@gmail.com | +91 88476 27161 · +91 97797 55748
arjunggullbadhar.com | linkedin.com/in/arjun-gullbadhar-11600b167

PROFESSIONAL SUMMARY

AI engineer with five years across data, backend, and applied-LLM systems, now specializing in **production LLM systems and their evaluation**. At PharmaEdge, builds the evaluation, retrieval, vision, and research pipelines behind a pharma competitive-intelligence product. Previously built a 41-tool production MCP server for an SEO platform and shipped 10 consumer iOS apps solo. Conference speaker on LLM testing in production.

EXPERIENCE

AI Engineer, PharmaEdge.ai (pharmaceutical competitive intelligence) Nov 2025 – Present

- **Built the flagship news-intelligence system end to end**, from ingestion to client-ranked delivery: pharma news lands every 30 minutes, is clustered into stories, and is ranked per client against tracked drugs, competitors, and therapeutic areas; raised story-clustering F1 from 0.62 to 0.855 on a hand-labeled golden set by combining embedding similarity, LLM-extracted entities, and entity-overlap scoring with a news-type taxonomy.
- **Built the conference-intelligence analyzer end to end** — a multimodal pipeline that turns conference posters, booth photos, and slide decks into structured competitive intelligence, routing each asset through adaptive OCR and vision paths (GPT-4.1/4o) and digitizing Kaplan–Meier survival curves into typed data; cut per-asset processing time 12–42× across asset types.
- **Established the evaluation layer for every LLM service**: versioned golden datasets with per-merge regression and cost gates that block any prompt or model change that regresses quality or raises cost, LLM-as-judge review calibrated against human labels, and drift detection; drove recurring output-formatting failures to zero with instructor + Pydantic structured outputs.
- **Built the investor-relations call intelligence system**: chunked Whisper ASR transcribes live earnings calls, biased by a domain vocabulary mined from each company’s previous three quarters of calls; client-tuned watchlists flag relevant moments to the right client as the call happens, and every call ships with a generated prep brief of questions and updates to listen for, plus summaries and extracted pipeline milestones.
- **Built a multi-stage deep-research workflow and RAG chat** over the proprietary corpus — pgvector retrieval with reciprocal-rank fusion and LLM reranking, plus retrieval-verification and citation checks so generated claims trace back to sources; structured extraction over ClinicalTrials.gov feeds the corpus typed trial records.
- **Built the observability pipeline for the LLM services**, logging every model call with tokens, latency, cost, and status, and tracking quality and spend across services.

AI Engineer (Contract), Welzin.ai (SEO content platform) Aug 2025 – Oct 2025

- **Designed and built a production MCP server exposing 41 tools** for content generation, scheduling, social publishing, and Search Console insights, used live by the client’s agent workflows; event-ID deduplication ensures retries never publish duplicate social posts.
- **Delivered three client-facing workflows end to end** across NestJS and Python services: a SERP-grounded SEO wizard, a server-rendered carousel studio, and a LinkedIn-style social composer.

Indie iOS Developer & Technical Writer*Mar 2024 – Nov 2025*

- **Shipped 10 consumer iOS apps solo in Flutter**, owning product design, LLM backends, App Store review, localization for 15+ storefront markets, and RevenueCat subscription billing; several were LLM-powered, including a calorie tracker built on chained model calls.
- **Curated a domain training dataset and fine-tuned Qwen** (Unsloth SFT) for domain-specific generation.

Data Analyst, Zebra Technologies (Antuit.ai)*May 2021 – Jan 2024*

- **Built pricing-analytics models for GlaxoSmithKline** covering 50+ markets; outputs fed directly into annual pricing strategy reviews with client stakeholders.
- **Developed demand-forecasting solutions** for retail brands including J.Crew and Madewell, informing inventory and promotion decisions.
- **Designed and operated PySpark ETL pipelines on Azure Databricks** processing 10M+ transactions daily.

TALKS & WRITING

- **Speaker, AAIF & CNCG Summit 2026**: “Prompt Regression Testing using Argo”
- **Speaker, OpenSearch Delhi Meetup 2026**: “Make Sense of a Pile of Stuff: Embeddings Beyond Search”
- **Medium Top Writer in Artificial Intelligence** (awarded Dec 2022) — 250K+ cumulative readers on the technical blog, medium.com/@arjungullbadhar

SKILLS

LLM & agents: OpenAI, Anthropic (Claude), and Gemini APIs; AWS Bedrock; Azure AI Foundry; structured outputs (instructor, Pydantic, PydanticAI); LangGraph, LangChain, LlamaIndex; MCP server design; multi-agent orchestration; LoRA/QLoRA fine-tuning

Evaluation & observability: golden-set regression testing, per-merge quality and cost gates, LLM-as-judge calibration, RAGAS, drift detection, LangSmith/Langfuse, full model-call logging

Retrieval & multimodal: PostgreSQL/pgvector, hybrid search, reciprocal-rank fusion, LLM reranking, OCR/vision extraction, Whisper ASR

Engineering & data: Python (async, typing), SQL, Bash; FastAPI, Docker, GitLab CI, SSE streaming; AWS (EC2, S3, CodePipeline); pandas, scikit-learn, PySpark, Azure Databricks; Flutter (iOS, full App Store lifecycle)

EDUCATION & AWARDS

B.E. Computer Science & Engineering (Cloud Computing with IBM)*2016 – 2020*

Chandigarh University

Kaggle Elo Merchant top 19% global · Hackoverflow Hackathon 2nd of 108 teams · TechGig ML Code Gladiators national finalist